



Lecture 5: Decision Trees, Random Forests and the Bias-Variance Tradeoff

Data Science Decal

Hosted by Machine Learning at Berkeley

Agenda

Background

Decision Tree Setup

Decision Tree Algorithm and Justification

Random Forests

The Bias-Variance Tradeoff

Questions

Background

The game is 20 Questions. Let's review the objective and rules:

The game is 20 Questions. Let's review the objective and rules:

- Player 1 thinks of a particular thing in a broad category e.g. People, Animals, Vegetables, etc.

The game is 20 Questions. Let's review the objective and rules:

- Player 1 thinks of a particular thing in a broad category e.g. People, Animals, Vegetables, etc.
- Player 2 must determine what Player 1 is thinking of.

The game is 20 Questions. Let's review the objective and rules:

- Player 1 thinks of a particular thing in a broad category e.g. People, Animals, Vegetables, etc.
- Player 2 must determine what Player 1 is thinking of.
- Player 2 can ask up to 20 **Yes or No** questions of Player 1, and he is able to correctly guess the thing, then Player 2 wins.

Alright. I'm Player 1, and you are going to be Player 2. I'm thinking of a Person. Which of the following questions are you most likely to **ask me first**? Why?

Alright. I'm Player 1, and you are going to be Player 2. I'm thinking of a Person. Which of the following questions are you most likely to **ask me first**? Why?

- Is the person of from Hawaii?



Alright. I'm Player 1, and you are going to be Player 2. I'm thinking of a Person. Which of the following questions are you most likely to **ask me first**? Why?

- Is the person of from Hawaii?
- Is your person male?



You probably chose "**Is your person male?**" Why is this obviously the best choice?

You probably chose "**Is your person male?**" Why is this obviously the best choice?

- You can automatically **eliminate half** of the candidates for Player 1's choice.

You probably chose "**Is your person male?**" Why is this obviously the best choice?

- You can automatically **eliminate half** of the candidates for Player 1's choice.
- Compared to asking the Hawaii question:

You probably chose "**Is your person male?**" Why is this obviously the best choice?

- You can automatically **eliminate half** of the candidates for Player 1's choice.
- Compared to asking the Hawaii question:
 - If from Hawaii: You have awesomely good luck. You'll likely win soon.

You probably chose "**Is your person male?**" Why is this obviously the best choice?

- You can automatically **eliminate half** of the candidates for Player 1's choice.
- Compared to asking the Hawaii question:
 - If from Hawaii: You have awesomely good luck. You'll likely win soon.
 - If not from Hawaii: You basically **wasted a question**.

Decision Tree Setup

- The **Decision Tree** is a model with a tree structure that encodes a collection of **if, then** statements in the **internal nodes**.

- The **Decision Tree** is a model with a tree structure that encodes a collection of **if, then** statements in the **internal nodes**.
- A query data point's features are used to channel the point to a **leaf node**, where it is assigned a label (classification) or value (regression).

- The **Decision Tree** is a model with a tree structure that encodes a collection of **if, then** statements in the **internal nodes**.
- A query data point's features are used to channel the point to a **leaf node**, where it is assigned a label (classification) or value (regression).
- Works naturally for continuous and categorical features!

It's like playing 20 questions with a data point $x^{(i)}$ like so:

- We wish to "guess" $y^{(i)} \in \{0, 1\}$ for classification and $y^{(i)} \in \mathbb{R}$ for regression.

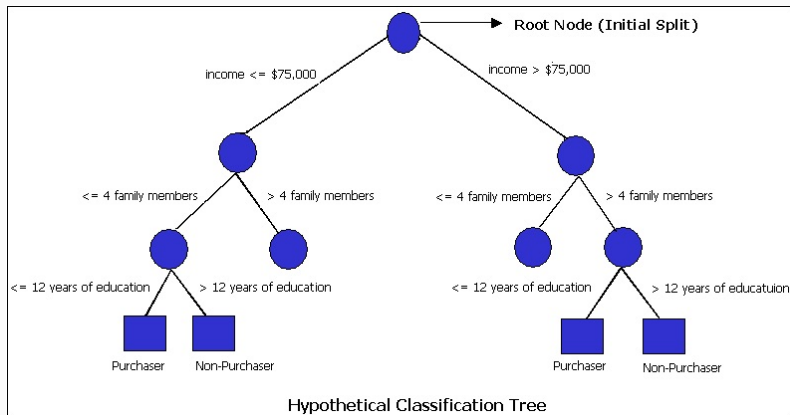
It's like playing 20 questions with a data point $x^{(i)}$ like so:

- We wish to "guess" $y^{(i)} \in \{0, 1\}$ for classification and $y^{(i)} \in \mathbb{R}$ for regression.
- We "ask" only binary "questions" of $x^{(i)}$, e.g. $x_3^{(i)} \geq 5?$.

It's like playing 20 questions with a data point $x^{(i)}$ like so:

- We wish to "guess" $y^{(i)} \in \{0, 1\}$ for classification and $y^{(i)} \in \mathbb{R}$ for regression.
- We "ask" only binary "questions" of $x^{(i)}$, e.g. $x_3^{(i)} \geq 5?$.
- These questions, or more formally **data splits** are what parameterize our model.

Decision Tree Example



Somehow, we must use our training dataset to learn the **features** and **values** to choose all of the splits shown.

Decision Tree Algorithm and Justification

We build the decision tree with the following **recursive** algorithm.

β is the value of the feature α on which the point will be split.

GrowTree(Set $S \subseteq \{1 \dots n\}$):

- if StopCriteria(S) is True:

We build the decision tree with the following **recursive** algorithm.

β is the value of the feature α on which the point will be split.

GrowTree(Set $S \subseteq \{1 \dots n\}$):

- if StopCriteria(S) is True:
 - return new *Leaf*(Value(S))

We build the decision tree with the following **recursive** algorithm.

β is the value of the feature α on which the point will be split.

GrowTree(Set $S \subseteq \{1 \dots n\}$):

- if StopCriteria(S) is True:
 - return new *Leaf*(Value(S))
- $\alpha, \beta \leftarrow \text{findSplit}(S)$

We build the decision tree with the following **recursive** algorithm.

β is the value of the feature α on which the point will be split.

GrowTree(Set $S \subseteq \{1 \dots n\}$):

- if StopCriteria(S) is True:
 - return new *Leaf*(Value(S))
- $\alpha, \beta \leftarrow \text{findSplit}(S)$
- $S_l \leftarrow \{i : x_\alpha^{(i)} \leq \beta\}$ // left split

We build the decision tree with the following **recursive** algorithm.

β is the value of the feature α on which the point will be split.

GrowTree(Set $S \subseteq \{1 \dots n\}$):

- if StopCriteria(S) is True:
 - return new *Leaf*(Value(S))
- $\alpha, \beta \leftarrow \text{findSplit}(S)$
- $S_l \leftarrow \{i : x_\alpha^{(i)} \leq \beta\}$ // left split
- $S_r \leftarrow \{i : x_\alpha^{(i)} > \beta\}$ // right split

We build the decision tree with the following **recursive** algorithm.

β is the value of the feature α on which the point will be split.

GrowTree(Set $S \subseteq \{1 \dots n\}$):

- if StopCriteria(S) is True:
 - return new *Leaf*(Value(S))
- $\alpha, \beta \leftarrow \text{findSplit}(S)$
- $S_l \leftarrow \{i : x_\alpha^{(i)} \leq \beta\}$ // left split
- $S_r \leftarrow \{i : x_\alpha^{(i)} > \beta\}$ // right split
- return new *InternalNode*(α, β , GrowTree(S_l), GrowTree(S_r))

We have shown how to find ideal parameters θ using **gradient descent**. This works because we are **optimizing a loss function**, $J(\theta)$.

Ideal split criteria in decision trees are chosen through minimizing the amount of **entropy** in each split group.

- In Data Science, entropy \rightarrow uncertainty \rightarrow surprise in knowing label

- In Data Science, entropy \rightarrow uncertainty \rightarrow surprise in knowing label
- $p_c = P(y_i = c) \rightarrow H(p_c) = -\log_2 p_c$


- In Data Science, entropy \rightarrow uncertainty \rightarrow surprise in knowing label
- $p_c = P(y_i = c) \rightarrow H(p_c) = -\log_2 p_c$
- What is $H(0)$? $H(1)$? $H(\frac{1}{2})$?

Data points belong to total of c classes, and fraction of data points in each is $[p_0 \dots p_{c-1}] : \sum_i p_i = 1$, then **entropy (H) of the set** is:

$$H = - \sum_{i=0}^{c-1} p_i \log p_i$$

(Using log base 2) If we have two classes, what values of p_0, p_1 will maximize this function? Do you see the connection to high entropy?

Here is a graph depicting $H(p_0)$ with $p_1 = 1 - p_0$:



To summarize: A Set S will have low entropy when $|p_0 - p_1|$ is large.

Think back to 20 Questions.

Think back to 20 Questions.

- The questions that give you the **most information gain** are the ones that **minimize entropy** in the two branches - S_l, S_r

Think back to 20 Questions.

- The questions that give you the **most information gain** are the ones that **minimize entropy** in the two branches - S_l, S_r
- We weight the size of the sets: $|S_l|, |S_r|$ to **discourage lopsided splits** unless necessary.

Think back to 20 Questions.

- The questions that give you the **most information gain** are the ones that **minimize entropy** in the two branches - S_l, S_r
- We weight the size of the sets: $|S_l|, |S_r|$ to **discourage lopsided splits** unless necessary.

Think back to 20 Questions.

- The questions that give you the **most information gain** are the ones that **minimize entropy** in the two branches - S_l, S_r
- We weight the size of the sets: $|S_l|, |S_r|$ to **discourage lopsided splits** unless necessary.


The **entropy of a split** is given as

$$H_{after} = \frac{|S_l|H_{S_l} + |S_r|H_{S_r}}{|S_l| + |S_r|}$$


Visualizing Optimal Split



$H(p_C)$ entropy: strictly concave



% misclassified: concave, not strict



Hence, $\text{findSplit}(S)$ gives split from:

$$\arg \min H_{\text{after}}(S_l, S_r)$$

Hence, $\text{findSplit}(S)$ gives split from:

$$\arg \min H_{\text{after}}(S_l, S_r)$$

There is no analytical solution for the problem above. Best option is checking each split (α, β) manually.


Hence, $\text{findSplit}(S)$ gives split from:

$$\arg \min H_{\text{after}}(S_l, S_r)$$

There is no analytical solution for the problem above. Best option is checking each split (α, β) manually.

- With n data points, d features, and up to k splits on a feature: runtime is $\mathcal{O}(ndk)$

Sorting the points once per feature, we can evaluate entropy for that feature on each split in $\mathcal{O}(1)$ time. This gives us an overall runtime of $\mathcal{O}(nd \log n)$




Finished tree can divide feature space into arbitrary number of *axis-aligned* regions.

Decision Trees vs. Linear Models – Depends on underlying data structure


1. True linear boundary:

Linear model fits perfectly




3. True linear boundary:

Flexible decision tree can approximate




2. True nonlinear boundary:

Linear model fits poorly



4. True nonlinear boundary:

Simple decision tree fits perfectly




Source: "An Introduction to Statistical Learning with Applications in R", James et al., 2013


© OLIVER WYMAN

24

Image Credits to Professor Shewchuk



Which one would you rather have? Points are **training data**.



Two common strategies:

Two common strategies:

- **Early Stopping:** Limit depth, threshold purity of leaves (% same class or set range for regression), few data points, tiny volume

Two common strategies:

- **Early Stopping:** Limit depth, threshold purity of leaves (% same class or set range for regression), few data points, tiny volume
- **Pruning:** Create whole tree. Greedily remove leaves, test for improved validation accuracy.

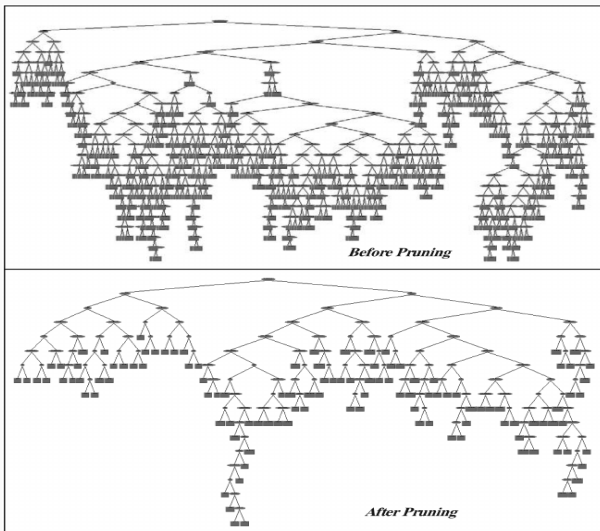
Two common strategies:

- **Early Stopping:** Limit depth, threshold purity of leaves (% same class or set range for regression), few data points, tiny volume
- **Pruning:** Create whole tree. Greedily remove leaves, test for improved validation accuracy.

Two common strategies:

- **Early Stopping:** Limit depth, threshold purity of leaves (% same class or set range for regression), few data points, tiny volume
- **Pruning:** Create whole tree. Greedily remove leaves, test for improved validation accuracy.

Impure leaves return majority vote (classification), average value (regression).



Let's see a cool application of Decision Trees to a somewhat-practical application: Click Me!



Random Forests

Back to 20 questions.

Back to 20 questions.

- If many people play against you, it is more likely that **at least one of them** figures it out than the individual chance someone figures it out.

Back to 20 questions.

- If many people play against you, it is more likely that **at least one of them** figures it out than the individual chance someone figures it out.
- Some people are better at asking questions than others.

Back to 20 questions.

- If many people play against you, it is more likely that **at least one of them** figures it out than the individual chance someone figures it out.
- Some people are better at asking questions than others.

Back to 20 questions.

- If many people play against you, it is more likely that **at least one of them** figures it out than the individual chance someone figures it out.
- Some people are better at asking questions than others.

What if we made one model by **combining many decision trees**?
Combine their judgments into something nuanced.

Original decision tree algorithm → **every tree would be identical**. Encourage diversity in finding splits:

Original decision tree algorithm → **every tree would be identical**. Encourage diversity in finding splits:

- Each tree gets **different data**. Sample points $x^{(i)}$ **with replacement** to form different starting sets S .

Original decision tree algorithm → **every tree would be identical**. Encourage diversity in finding splits:

- Each tree gets **different data**. Sample points $x^{(i)}$ **with replacement** to form different starting sets S .
- At each call of `GrowTree()`, constrict the `FindSplit()` method to examine only a **subset of features** to split on.

Original decision tree algorithm → **every tree would be identical**. Encourage diversity in finding splits:

- Each tree gets **different data**. Sample points $x^{(i)}$ **with replacement** to form different starting sets S .
- At each call of `GrowTree()`, constrict the `FindSplit()` method to examine only a **subset of features** to split on.


Original decision tree algorithm → **every tree would be identical**. Encourage diversity in finding splits:

- Each tree gets **different data**. Sample points $x^{(i)}$ **with replacement** to form different starting sets S .
- At each call of `GrowTree()`, constrict the `FindSplit()` method to examine only a **subset of features** to split on.

These techniques are called **bootstrap aggregating (bagging)** and **random feature selection** respectively.

How Much Does This Cow Weigh?

(All People)



Why do the aforementioned techniques work?

Why do the aforementioned techniques work?

- **Bagging:** Points (almost surely) repeated in data \rightarrow weighted more \rightarrow minimizes the effect of outliers.

Why do the aforementioned techniques work?

- **Bagging:** Points (almost surely) repeated in data \rightarrow weighted more \rightarrow minimizes the effect of outliers.
- **Random feature selection:** Despite bagging, trees may still split on same first feature. This forces different "perspectives" on split paths.

Why do the aforementioned techniques work?

- **Bagging:** Points (almost surely) repeated in data \rightarrow weighted more \rightarrow minimizes the effect of outliers.
- **Random feature selection:** Despite bagging, trees may still split on same first feature. This forces different "perspectives" on split paths.
 - *Example:* Asking "Is the person still alive?" might be better than asking if the person is male, *depending on the situation*.

The aforementioned techniques also work for other models.

The aforementioned techniques also work for other models.

- **Bagging:** for SVM's, logistic regression, # of times point is picked \rightarrow point's weight in loss function. It results in $\approx 62.3\%$ of data represented.

The aforementioned techniques also work for other models.

- **Bagging:** for SVM's, logistic regression, # of times point is picked \rightarrow point's weight in loss function. It results in $\approx 62.3\%$ of data represented.
- **Random Feature Selection** \rightarrow A way to address high dimensionality on optimization based strategies.

The aforementioned techniques also work for other models.

- **Bagging:** for SVM's, logistic regression, # of times point is picked \rightarrow point's weight in loss function. It results in $\approx 62.3\%$ of data represented.
- **Random Feature Selection** \rightarrow A way to address high dimensionality on optimization based strategies.
 - For random forests, \sqrt{d} , $\frac{d}{3}$ features are recommended for splits on classification and regression respectively.

A **random forest** is defined to be:

A **random forest** is defined to be:

- One model

A **random forest** is defined to be:

- One model
- Composed of multiple decision trees

A **random forest** is defined to be:

- One model
- Composed of multiple decision trees
- Which all use both, bagging and random feature selection


A **random forest** is defined to be:

- One model
- Composed of multiple decision trees
- Which all use both, bagging and random feature selection
- And work by casting majority votes (classification) or averages (regression), with values made from all of the constituent decision trees.

Greatest advantage of random forests is their ability to make **arbitrarily complex decision boundaries** given a large enough number of trees.

Greatest advantage of random forests is their ability to make **arbitrarily complex decision boundaries** given a large enough number of trees.

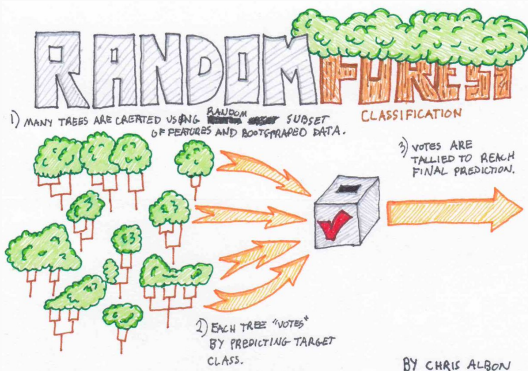
Despite this, random forests are less likely to overfit.



Top row: depth = 4, Bottom row: depth = 12.


Left: 1 random feature per split, Middle: 5 features, Right: 50 features

Let's check out how our favorite data science model library presents Random Forests: Click Me!



The Bias-Variance Tradeoff

Overfitting and underfitting



We start by assuming our data comes from a true function f but also has random noise ϵ added afterward. We have a hypothesis h which is always imperfect for 2 reasons:

- **Bias:** h is unable to fit f perfectly because it lacks capacity or complexity. (Fitting a line to a parabola).

We start by assuming our data comes from a true function f but also has random noise ϵ added afterward. We have a hypothesis h which is always imperfect for 2 reasons:

- **Bias:** h is unable to fit f perfectly because it lacks capacity or complexity. (Fitting a line to a parabola).
- **Variance:** h is fitting to the noise in the data, thus missing the true function f .

We start by assuming our data comes from a true function f but also has random noise ϵ added afterward. We have a hypothesis h which is always imperfect for 2 reasons:

- **Bias:** h is unable to fit f perfectly because it lacks capacity or complexity. (Fitting a line to a parabola).
- **Variance:** h is fitting to the noise in the data, thus missing the true function f .

We start by assuming our data comes from a true function f but also has random noise ϵ added afterward. We have a hypothesis h which is always imperfect for 2 reasons:

- **Bias:** h is unable to fit f perfectly because it lacks capacity or complexity. (Fitting a line to a parabola).
- **Variance:** h is fitting to the noise in the data, thus missing the true function f .

Now, suppose we have an arbitrary point z , and a generated data point $\gamma = f(z) + \epsilon$.

Since ϵ is a random variable, γ is now one as well. This means we can calculate statistics such as expected value and variance.

Since ϵ is a random variable, γ is now one as well. This means we can calculate statistics such as expected value and variance.

- $\mathbb{E}[\gamma] = f(z)$ and $\text{Var}[\gamma] = \text{Var}[\epsilon]$.

Since ϵ is a random variable, γ is now one as well. This means we can calculate statistics such as expected value and variance.

- $\mathbb{E}[\gamma] = f(z)$ and $\text{Var}[\gamma] = \text{Var}[\epsilon]$.

Since ϵ is a random variable, γ is now one as well. This means we can calculate statistics such as expected value and variance.

- $\mathbb{E}[\gamma] = f(z)$ and $\text{Var}[\gamma] = \text{Var}[\epsilon]$.

Let's analyze the Expectation of Loss when we use Mean-Squared Error as a cost function:

Since ϵ is a random variable, γ is now one as well. This means we can calculate statistics such as expected value and variance.

- $\mathbb{E}[\gamma] = f(z)$ and $\text{Var}[\gamma] = \text{Var}[\epsilon]$.

Let's analyze the Expectation of Loss when we use Mean-Squared Error as a cost function:

$$\begin{aligned} R(h) &= \mathbb{E}[(h(z) - \gamma)^2] \\ &= \mathbb{E}[h(z)^2] + \mathbb{E}[\gamma^2] - 2\mathbb{E}[\gamma h(z)] \end{aligned}$$

- We note that γ and $h(z)$ are independent. This allows us to exploit the property: $\mathbb{E}[XY] = \mathbb{E}[X]\mathbb{E}[Y]$.

- We note that γ and $h(z)$ are independent. This allows us to exploit the property: $\mathbb{E}[XY] = \mathbb{E}[X]\mathbb{E}[Y]$.
- Since $\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$, we can substitute $\mathbb{E}[X^2] = \text{Var}[X] + \mathbb{E}[X]^2$ in the derivation

- We note that γ and $h(z)$ are independent. This allows us to exploit the property: $\mathbb{E}[XY] = \mathbb{E}[X]\mathbb{E}[Y]$.
- Since $\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$, we can substitute $\mathbb{E}[X^2] = \text{Var}[X] + \mathbb{E}[X]^2$ in the derivation

- We note that γ and $h(z)$ are independent. This allows us to exploit the property: $\mathbb{E}[XY] = \mathbb{E}[X]\mathbb{E}[Y]$.
- Since $\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$, we can substitute $\mathbb{E}[X^2] = \text{Var}[X] + \mathbb{E}[X]^2$ in the derivation

$$\mathbb{E}[h(z)^2] + \mathbb{E}[\gamma^2] - 2\mathbb{E}[\gamma h(z)]$$

- We note that γ and $h(z)$ are independent. This allows us to exploit the property: $\mathbb{E}[XY] = \mathbb{E}[X]\mathbb{E}[Y]$.
- Since $\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$, we can substitute $\mathbb{E}[X^2] = \text{Var}[X] + \mathbb{E}[X]^2$ in the derivation

$$\mathbb{E}[h(z)^2] + \mathbb{E}[\gamma^2] - 2\mathbb{E}[\gamma h(z)]$$

$$\begin{aligned} &= \text{Var}[h(z)] + \mathbb{E}[h(z)]^2 + \text{Var}[\gamma] + \mathbb{E}[\gamma]^2 - 2\mathbb{E}[\gamma]\mathbb{E}[h(z)] \\ &= (\mathbb{E}[h(z)] - \mathbb{E}[\gamma])^2 + \text{Var}[h(z)] + \text{Var}[\gamma] \\ &= (\mathbb{E}[h(z)] - f(z))^2 + \text{Var}[h(z)] + \text{Var}[\epsilon] \end{aligned}$$

From the last slide, we see three important quantities:

From the last slide, we see three important quantities:

- $(\mathbb{E}[h(z)] - f(z))^2$ is the squared **bias**. It is how much the model expects to differ from the real data distribution function.

From the last slide, we see three important quantities:

- $(\mathbb{E}[h(z)] - f(z))^2$ is the squared **bias**. It is how much the model expects to differ from the real data distribution function.
- $\text{Var}[h(z)]$ is the **variance**. The risk is explicitly composed partly by the variance of our model.

From the last slide, we see three important quantities:

- $(\mathbb{E}[h(z)] - f(z))^2$ is the squared **bias**. It is how much the model expects to differ from the real data distribution function.
- $\text{Var}[h(z)]$ is the **variance**. The risk is explicitly composed partly by the variance of our model.
- $\text{Var}[\epsilon]$ is the **irreducible error**. The noise put into the data is entirely out of our control.

From the last slide, we see three important quantities:


- $(\mathbb{E}[h(z)] - f(z))^2$ is the squared **bias**. It is how much the model expects to differ from the real data distribution function.
- $\text{Var}[h(z)]$ is the **variance**. The risk is explicitly composed partly by the variance of our model.
- $\text{Var}[\epsilon]$ is the **irreducible error**. The noise put into the data is entirely out of our control.

From the last slide, we see three important quantities:

- $(\mathbb{E}[h(z)] - f(z))^2$ is the squared **bias**. It is how much the model expects to differ from the real data distribution function.
- $\text{Var}[h(z)]$ is the **variance**. The risk is explicitly composed partly by the variance of our model.
- $\text{Var}[\epsilon]$ is the **irreducible error**. The noise put into the data is entirely out of our control.


Because we are able to break down the risk of any hypothesis like so, this process is called the **Bias-Variance Decomposition**.

Visualizing Each of the Quantities



50 fits (20 examples each)

=



Lots of issues and techniques center around adjusting or accounting for bias, variance and irreducible error.

Lots of issues and techniques center around adjusting or accounting for bias, variance and irreducible error.

- Underfitting directly relates to having too much bias.

Lots of issues and techniques center around adjusting or accounting for bias, variance and irreducible error.

- Underfitting directly relates to having too much bias.
- Overfitting is most often a result of too much variance.

Lots of issues and techniques center around adjusting or accounting for bias, variance and irreducible error.

- Underfitting directly relates to having too much bias.
- Overfitting is most often a result of too much variance.
- Training error shows bias, not variance. Test error reflects both.

Lots of issues and techniques center around adjusting or accounting for bias, variance and irreducible error.

- Underfitting directly relates to having too much bias.
- Overfitting is most often a result of too much variance.
- Training error shows bias, not variance. Test error reflects both.
- For many distributions, variance $\rightarrow 0$ as $n \rightarrow \infty$.

Lots of issues and techniques center around adjusting or accounting for bias, variance and irreducible error.

- Underfitting directly relates to having too much bias.
- Overfitting is most often a result of too much variance.
- Training error shows bias, not variance. Test error reflects both.
- For many distributions, variance $\rightarrow 0$ as $n \rightarrow \infty$.
- Assuming h has sufficient modeling capacity, bias $\rightarrow 0$ as $n \rightarrow \infty$


- **Adding features:** higher quality feature \rightarrow higher drop in bias.

- **Adding features:** higher quality feature \rightarrow higher drop in bias.
- *Variance always increases.* \rightarrow only add feature if bias drop variance gain.

- **Adding features:** higher quality feature \rightarrow higher drop in bias.
- *Variance always increases.* \rightarrow only add feature if bias drop variance gain.
- Noise in test set only affects $Var(\epsilon)$. Noise in training set affects only bias and $Var(\epsilon)$.

- **Adding features:** higher quality feature \rightarrow higher drop in bias.
- *Variance always increases.* \rightarrow only add feature if bias drop variance gain.
- Noise in test set only affects $Var(\epsilon)$. Noise in training set affects only bias and $Var(\epsilon)$.
- While f is unknowable, we can pick \hat{f} , synthesize data and test models for strengths and weaknesses.

Balancing Bias and Variance



We decision trees to linear and logistic regression, SVM's.

- **Bias:** Decision trees have highly nonlinear decision boundaries and its parameters can be arbitrarily complex.

We compare decision trees to linear and logistic regression, SVM's.

- **Bias:** Decision trees have highly nonlinear decision boundaries and its parameters can be arbitrarily complex.
 - Decision trees have relatively **low bias**.

We decision trees to linear and logistic regression, SVM's.

- **Bias:** Decision trees have highly nonlinear decision boundaries and its parameters can be arbitrarily complex.
 - Decision trees have relatively **low bias**.
- **Variance:** Decision tree leaves will be pure (without early stopping/pruning), so all noise is accounted for. Other regression models have clear regularization.

We decision trees to linear and logistic regression, SVM's.

- **Bias:** Decision trees have highly nonlinear decision boundaries and its parameters can be arbitrarily complex.
 - Decision trees have relatively **low bias**.
- **Variance:** Decision tree leaves will be pure (without early stopping/pruning), so all noise is accounted for. Other regression models have clear regularization.
 - Decision trees have relatively **high variance**.

- **Bias:** Random Forests average decision tree decision boundaries, so they are no longer axis-aligned.

- **Bias:** Random Forests average decision tree decision boundaries, so they are no longer axis-aligned.
 - Random Forests have **lower bias.** than decision trees.

- **Bias:** Random Forests average decision tree decision boundaries, so they are no longer axis-aligned.
 - Random Forests have **lower bias.** than decision trees.
- **Variance:** Random Forests average out decision trees.

$$\text{Var}\left[\frac{1}{n}(h_1(z) + \dots h_n(z))\right] = \frac{1}{n^2}n\text{Var}[h_i(z)] = \frac{1}{n}\text{Var}[h_i(z)]$$

- **Bias:** Random Forests average decision tree decision boundaries, so they are no longer axis-aligned.
 - Random Forests have **lower bias.** than decision trees.
- **Variance:** Random Forests average out decision trees.

$$\text{Var}\left[\frac{1}{n}(h_1(z) + \dots h_n(z))\right] = \frac{1}{n^2}n\text{Var}[h_i(z)] = \frac{1}{n}\text{Var}[h_i(z)]$$

- Because of the properties of Variance, Random Forests have **lower variance** than decision trees.

Especially because of variance, an **ensemble learner** comprised of several models that vote or average will generalize better.


- The main consideration simply becomes training costs

Especially because of variance, an **ensemble learner** comprised of several models that vote or average will generalize better.


- The main consideration simply becomes training costs
- Prior to deep learning revival in 2010's, random forests and ensembled SVM's were state-of-the-art for data science classification, prediction.

Especially because of variance, an **ensemble learner** comprised of several models that vote or average will generalize better.

- The main consideration simply becomes training costs
- Prior to deep learning revival in 2010's, random forests and ensembled SVM's were state-of-the-art for data science classification, prediction.
- Even now, can ensemble deep learning models e.g. dropout, mixture of experts




Training
 $p_{\text{keep}} = 0.75$



Test
 $p_{\text{keep}} = 1.0$

Bonus: Mixture of Experts



Questions

Questions?